

- 將中斷點設定在「syscall_init」，並回答下列問題

1. 當應用程式發出system call時，Linux會從哪裡開始執行Linux kernel

從start_kernel 開始執行trap_init，再從trap_init 執行cpu_init，cpu_init執行syscall_init。

```

Reading symbols from /home/parallels/osdi2019-master/kernel-4.0/vmlinux...done.
(gdb) b start_kernel
Breakpoint 1 at 0xffffffff822f5f20: file init/main.c, line 490.
(gdb) c
Continuing.

Breakpoint 1, start_kernel () at init/main.c:490
490 {
(gdb) b syscall_init
Breakpoint 2 at 0xffffffff81025082: file arch/x86/kernel/cpu/common.c, line 1172.
(gdb) c
Continuing.

Breakpoint 2, syscall_init () at arch/x86/kernel/cpu/common.c:1172
1172 wrmsrl(MSR_STAR, ((u64)_USER32_CS)<<48 | ((u64)_KERNEL_CS)<<32);
(gdb) #3  0xffffffff822f6093 in start_kernel () at init/main.c:550
550     trap_init();
(gdb)

```

▼ Thread #1 1 (CPU#0 [running]) (Suspended : Breakpoint)

- ≡ syscall_init() at common.c:1,172 0xffffffff81025082
- ≡ cpu_init() at common.c:1,336 0xffffffff81025082
- ≡ trap_init() at traps.c:1,006 0xffffffff822f9b82
- ≡ start_kernel() at main.c:550 0xffffffff822f6093
- ≡ x86_64_start_reservations() at head64.c:200 0xffffffff822f6093
- ≡ x86_64_start_kernel() at head64.c:189 0xffffffff822f6093

start_kernel:

```

547     pidhash_init();
548     vfs_caches_init_early();
549     sort_main_extable();
550     trap_init();
551     mm_init();
552

```

trap_init:

```

1003     /*
1004      * Should be a barrier for any external CPU
1005      */
1006     cpu_init();
1007
1008     x86_init.irqs.trap_init();
1009

```

cpu_init:

```

1334
1335     memset(me->thread.tls_array, 0, GDT_ENTRY_TLS_ENTRIES * 8);
1336     syscall_init();
1337
1338     wrmsrl(MSR_FS_BASE, 0);

```

syscall_init:

```

1171     /*
1172     * Set the system call gate to the kernel code.
1173     * The gate is a 32-bit instruction, so we need to
1174     * use the 32-bit instruction format.
1175     */
1176     wrmsrl(MSR_STAR, ((u64)_USER32_CS)<<48 | ((u64)_KERNEL_CS)<<32);
1177     wrmsrl(MSR_LSTAR, system_call);
1178     wrmsrl(MSR_CSTAR, ignore_sysret);

```

2. 如何使用rax暫存器呼叫對應的system call處理函數？

使用sys_call_table可以找到對應的處理函數

```
(gdb) b system_call
Breakpoint 3 at 0xffffffff81d5d600: file arch/x86/kernel/entry_64.S, line 330.
(gdb) c
Continuing.

Breakpoint 3, <signal handler called>
(gdb) si
<signal handler called>
```

```
ffffffff81d5d57f:    movq %r10,%rcx
ffffffff81d5d588:    mov    %r10,%rcx
267          call *sys_call_table(,%rax,8) # XXX: rdi
          callq *-0x7e1fa140(,%rax,8)
```

透過gdb 知道是63號 system call:

```
<signal handler called>
(gdb) info registers rax
rax          0x3f      63
```

第63號為newuname:

```
ffffffff81d5dae0 <stub_execve>, 0xffffffff81092085 <Sys_exit>, 0xffffffff81093d8a <Sys_wait4>, 0xffffffff810a7263 <
ill>, 0xffffffff810ac660 <Sys_newuname>, 0xffffffff8148e8f4 <Sys_semget>, 0xffffffff81491417 <Sys_semop>, 0xfffff
490731 <Sys_semctl>, 0xffffffff8149473f <Sys_shmdt>, 0xffffffff8148bbca <Sys_msqget>, 0xffffffff8148caa5 <Sys_msc
```